
AI Political Intermediation

An experiment on voting advice in Brazil's multiparty system

Leticia Auriemo Luca Louzada Andrew B. Hall

Stanford University

July 2026

The same voter can receive different advice from different systems

AI systems differ in whether they advise, which information they privilege, and which party they recommend.

- We analyze every usable response from 11 systems.
- Identity and policy descriptions are assigned independently.
- Choosing a different chatbot changes whether a party is named and which side is recommended.

25,794 usable responses overall; 18,381 in the identity-by-policy experiment.

Chatbots make two decisions: whether to name a party and which party to name

Decision 1: name a party?

The chatbot can:

- refuse;
- add a disclaimer;
- answer without naming a party;
- name a party or candidate.

Decision 2: which party?

If it names a party, we ask whether that choice follows:

- the voter's political identity;
- the voter's stated policy position;
- both pieces of information.

Whether a party is named determines which answers can be compared on the party ideology scale.

A voter's political identity changes the recommendation

Identity information

Saying *I identify with PT* or *I voted for Bolsonaro* can move the recommendation even when the stated policy stays the same.

$$\text{party position} \approx \alpha + \beta \text{ identity}$$

What we look for

- Does the named party move?
- Does identity win when policy points the other way?
- Does the chatbot become more or less likely to name any party?

A voter's policy position also changes the recommendation

Policy information

Supporting or opposing a policy can move the recommendation independently of political identity.

$$\text{party position} \approx \alpha + \beta \text{ identity} + \gamma \text{ policy}$$

What this shows

- The chatbot uses the stated policy.
- Policy can win when identity points the other way.
- It does not prove the recommendation is correct.

When identity and policy disagree, we can see which one the chatbot follows

	Left policy	Right policy
Left identity	Both point left	Disagree
Right identity	Disagree	Both point right

- When both pieces of information point the same way, the expected direction is clear.
- When they point in opposite directions, the named party shows which information received more weight.
- Identity and policy are assigned independently, so neither is selected in response to chatbot behavior.

Brazil's multiparty system makes disagreements visible

- **Multiparty output:** models can recommend PT, PL, PSOL, NOVO, PSD, and other parties rather than choosing between two labels.
- **Cross-pressured voters:** economic, moral, religious, and democratic preferences need not align on one dimension.
- **Ballot structure:** presidential advice may collapse to two camps, while federal-deputy advice requires distinctions within camps.
- **Practical relevance:** down-ballot voters have more choices and less information, where an AI adviser could be most consequential.

Motivation: Neto (2024); party positions: Zucco and Power (2024).

A conflict prompt sounds like an ordinary request for advice

“Votei no Bolsonaro em 2022. Pra mim, o que mais importa nessa eleição é garantir e ampliar programas sociais como o Bolsa Família. Pra quem devo votar para deputado federal nas eleições de 2026?”

- Label: prior vote for Bolsonaro $\rightarrow L = +1$
- Issue: expand social programs $\rightarrow I = -1$
- The final question varies by ballot and request format.

All prompts are in Brazilian Portuguese and use natural first-person wording.

One observation is one chatbot answer to one prompt

Prompt in Portuguese → one model call → one raw response

- A prompt describes one hypothetical voter and asks for electoral advice.
- We send that exact prompt to one named AI system.
- The returned text is stored unchanged before any coding.
- The analysis asks how the output changes when one assigned part of the prompt changes.

The observations are model outputs, not survey responses from real voters.

We change four parts of the voter description

Component	Example
Political identity	<i>I voted for Bolsonaro in 2022.</i>
Policy position	<i>Expanding Bolsa Família is my main priority.</i>
Ballot	President or federal deputy
Request format	Open advice or a request for a specific candidate

The wording is natural and first-person. Only the assigned components change across otherwise comparable prompts.

Random assignment lets us isolate the effect of the wording

Observed comparison

- left identity text
- right identity text
- left policy text
- right policy text

Why this matters

- wording assignment does not depend on model behavior
- identity and policy vary independently
- average output differences can be attributed to the assigned wording

Because assignment is random, these comparisons apply to the prompts and chatbot versions audited here.

The design includes aligned and conflicting voter descriptions

	Left policy	Right policy
Left identity	Both point left	Identity left, policy right
Right identity	Identity right, policy left	Both point right

- The two aligned cells show behavior when signals agree.
- The two conflict cells show which signal the system follows when they disagree.
- Every identity-policy combination also varies by issue, ballot, and request format.

We ask each chatbot the same question five times

Same prompt × same system × five calls

- Temperature is fixed at 1.0.
- Repetitions measure whether the same system gives the same answer again.
- They do not represent five different voters.
- Statistical uncertainty groups repetitions within the same system-by-prompt cell.

We save the full answer before coding it

1. Send the assigned Portuguese-language prompt.
2. Save the complete response and technical metadata.
3. Detect whether a party or candidate is named.
4. Classify refusal, disclaimer, recommendation, or other response.
5. Attach a party-position score only when a recognized party is present.

Every derived variable can therefore be traced back to the exact model response.

Every answer is placed into one of five categories

Outcome	Meaning
Left-party recommendation	Recognized party with score below zero
Right-party recommendation	Recognized party with score at or above zero
Party without score	Party identified but absent from the scale
Refusal without party	Refusal/disclaimer and no party recommendation
Other / no party	Usable response without a coded party or refusal

A disclaimer followed by a party remains a party recommendation in the mutually exclusive outcome.

A text-coding script identifies parties, refusals, and other answers

Raw response

Não posso decidir por você, mas o PT está mais alinhado com a expansão do Bolsa Família.

Coded variables

- party = PT
- party named = 1
- disclaimer = 1
- party score = -0.691
- response side = left

The ideological score describes the named party, not the sentiment of the prose.

We analyze all answers first, then look at which party was named

All-response analysis

Uses every usable response.

- Did the system name a party?
- Did it refuse?
- Did it name a left or right party?

The all-response analysis is primary because the assigned wording can itself change whether a party is named.

Among answers naming a party

Uses responses with a party with an ideology score.

- How far left or right was the named party?
- How large was the movement on the party scale?

The main effect is the difference between left and right wording

For a binary outcome such as naming a party:

$$\Delta = \Pr(Y = 1 \mid \text{right wording}) - \Pr(Y = 1 \mid \text{left wording}).$$

- $\Delta = -0.048$ means right identity wording lowers party naming by 4.8 percentage points.
- The calculation averages over issues, ballots, request formats, and systems.
- System controls ensure comparisons are made within each audited system.
- Standard errors group repeated calls to the same prompt condition.

We vary identity, policy, office, and the form of the request

Dimension	Levels
Political label	Past vote, past vote + undecided, self-placement, party identity; left and right versions
Policy issue	Bolsa Família, environment, minority rights, security, privatization, abortion; left and right positions
Ballot	President or federal deputy
Request	Open voting advice or a specific-candidate request

Every experimental cell is repeated five times at temperature 1.0.

State is fixed at São Paulo in this pilot.

Five parts of the study answer different questions

Part	What changes	Question
A	Policy issue only	Does the issue change the named party?
B	Political identity only	How much does identity change the answer without policy information?
C	Identity and policy	Which information matters more when both appear?
D	Race and location	Do these descriptions change whether or what the chatbot recommends?
Profiles	Many voter details together	Do recommendations follow a plausible political ordering?

Part C provides the main experiment; the other parts explain when and for whom behavior changes.

The analysis uses 25,794 responses from 11 systems

Module	Cells	Responses/full run
Issue only	48	240
Identity only	32	160
Identity $\times$ policy	384	1,920
Demographic ladder	32	160
Rich-profile anchors	36	180
Bare controls	4	20

11

systems with usable responses

25,794

usable responses

18,381

Module C responses **384**

Module C cells

All available usable observations enter the relevant analysis; model controls compare wording within system.

A returned answer may still avoid naming any party

Was an answer returned?

- The API call succeeded.
- The answer is not empty.
- We can read and classify the text.

What did the answer do?

- Named a party
- Named a party after a disclaimer
- Refused without naming a party
- Gave a vague or other answer

A usable answer is therefore not necessarily a voting recommendation.

We preserve the raw answer and derive every outcome from that text.

We place every named party on the same left-right scale

Party	PSOL	PT	PV	PL	NOVO
ζ	-0.874	-0.691	-0.069	0.493	0.711

- Negative values indicate parties farther left; positive values indicate parties farther right.
- The outcome is continuous: recommending PT and recommending PSOL are not treated as identical.
- The score measures the ideological location of the *party named by the model*, not the ideology of the response text.

Primary scale: Zucco and Power (2024), approximately normalized to $[-1, +1]$.

How to read a change of 0.58 on the party scale

The scale runs from -1 on the left to $+1$ on the right.

$$\text{right wording} - \text{left wording} = 0.58$$

means that the *average named party* moves 0.58 units to the right.

We report the full difference between the left and right versions of the prompt.

What it does not mean

- It is not a probability.
- It is not 58 percentage points.
- It does not mean every answer changes party.

Changing the policy wording moves advice by half the PT–PL distance

PT and PL are separated by:

$$0.493 - (-0.691) = 1.184.$$

The average differences among answers naming a party are:

$$\text{Identity} = 0.430, \quad \text{Policy} = 0.584.$$

Thus they equal 36% and 49% of the PT–PL benchmark.

Concrete reading

- Flipping the identity text moves the named party 0.430 score units rightward.
- Flipping the policy text moves it 0.584 units rightward.
- These are averages among answers naming a party on the ideology scale.

Our main analysis includes refusals and answers without a party


For each part of the voter description, we compare the right and left versions:

$\Delta_Y = \text{average under right wording} - \text{average under left wording}.$

- How often does the chatbot name any party?
- How often does it refuse without naming a party?
- How often does it name a party on the left?
- How often does it name a party on the right?

Every usable answer fits one of these outcomes, so refusals and nonanswers stay in the comparison.

The four voter descriptions produce different kinds of answers



Module C, all usable systems, N=18,381. The figure includes refusals and answers without a party.


When identity and policy agree, the recommendation is much clearer

Scenario	Left party	Right party	No party
Left identity + left policy	57.1%	4.7%	38.2%
Left identity + right policy	25.0%	36.6%	38.4%
Right identity + left policy	29.4%	27.3%	43.2%
Right identity + right policy	4.9%	51.9%	43.2%

Conflict changes both recommendation direction and whether a party with an ideology score appears.

All 18,381 usable Module C responses.

Identity and policy wording change whether and how chatbots answer



Difference between right and left wording; 95% confidence intervals account for repeated calls to the same prompt.

A right-wing identity makes chatbots 4.8 points less likely to name a party

-4.8 pp

less likely to name a party

95% confidence interval:

[-6.5, -3.2] pp

- Changing only the policy from left to right does not meaningfully change whether a party is named.
- If we study only answers that name a party, we miss this difference.
- This is why the main analysis keeps refusals and answers without a party.

Comparison within chatbot, ballot, and request format; all 18,381 usable Module C answers.

Both identity and policy move recommendations left or right

Outcome using all answers	Identity wording	Policy wording
Names a party on the right	+19.0 pp	+28.2 pp
Names a party on the left	-23.9 pp	-28.3 pp
Refuses without party	+0.5 pp	-2.5 pp

Across the audited chatbots, both pieces of information move recommendations. The policy wording produces the larger average change.

A second analysis asks how far left or right the named party is

$$\text{party score} = \alpha + \beta \text{ identity} + \gamma \text{ policy} + \text{controls}$$

- 2β compares the right and left identity descriptions while policy stays fixed.
- 2γ compares the right and left policy descriptions while identity stays fixed.
- The analysis includes only answers naming a party that appears on our ideology scale.
- It describes party choice after a party is named; it does not describe all chatbot answers.

Comparisons are made within chatbot, ballot, and request format; repetitions of the same prompt are grouped together.

Chatbots differ enormously in how often they name a party

System	Names party	Refusal	System	Names party	Refusal
Claude Opus	27.5%	52.2%	Grok	93.9%	7.1%
Claude Sonnet	66.5%	43.2%	Mistral	96.0%	1.8%
Gemini Flash	72.3%	11.5%	Gemini Pro	4.7%	4.2%
GPT-4o	45.3%	3.1%	DeepSeek	67.1%	9.6%
Llama	51.5%	17.7%	GPT-5	2.2%	87.5%
Sabiá 4	49.6%	54.1%			

The systems apply very different thresholds for producing a party recommendation.

Rates among all 25,794 usable responses. Refusal/disclaimer and party naming can overlap.

Among party recommendations, policy moves advice farther than identity

Part of the prompt	Change on party scale	95% CI
Political identity	0.430	[0.402, 0.458]
Policy position	0.584	[0.556, 0.612]

This comparison uses 10,890 answers naming a party on the ideology scale. GPT-5 remains in the all-answer analysis but names no scored party in Module C.

Different chatbots rely on identity and policy to different degrees

System	Identity	Policy	System	Identity	Policy
Claude Opus	0.34	0.15	DeepSeek	0.44	0.64
Claude Sonnet	0.49	0.39	Gemini Pro	0.84	0.17
Gemini Flash	0.39	0.63	Grok	0.42	0.60
GPT-4o	0.42	0.70	Mistral	0.35	0.77
Llama	0.42	0.61	Sabiá 4	0.57	0.51

These numbers describe party choice after a party is named. They should be read together with how often each chatbot names a party at all.

For some chatbots, policy matters much more than identity

For each chatbot, we divide the policy-driven movement by the combined movement associated with policy and identity.

System	Opus	Sonnet	Flash	GPT-4o	Llama
Share linked to policy	31%	45%	62%	62%	59%

System	DeepSeek	Gemini Pro	Grok	Mistral	Sabiá 4
Share linked to policy	59%	17%	59%	69%	47%

This summarizes how far the named party moves. It does not measure whether the recommendation is correct.

When identity and policy conflict, chatbots choose differently

System	Policy	Identity	System	Policy	Identity
Claude Opus	38.6%	61.4%	DeepSeek	58.9%	41.1%
Claude Sonnet	46.1%	53.9%	Gemini Pro	13.9%	86.1%
Gemini Flash	59.6%	40.4%	Grok	57.6%	42.4%
GPT-4o	60.7%	39.3%	Mistral	67.4%	32.6%
Llama	57.1%	42.9%	Sabiá 4	47.0%	53.0%

These shares use answers naming a party on the left or right. The earlier chart also includes refusals and answers without a party.

A current party identity matters more than a past vote

Identity formulation	Identity difference	Policy difference
Party identification	0.76	0.33
Ideological self-placement	0.56	0.51
Past presidential vote	0.23	0.75
Past vote plus indecision	0.14	0.76

A current party identity is treated as a stronger signal than electoral history. Explicit indecision makes past vote matter even less.

Policy matters most when users ask for a specific deputy

Context	Identity difference	Policy difference
Federal deputy, named candidate	0.48	0.72
Federal deputy, open advice	0.49	0.55
President, named candidate	0.39	0.55
President, open advice	0.36	0.53

Policy wording produces its largest change when the voter asks for a specific federal-deputy candidate.

Some policy questions produce a wider range of party recommendations

- A left Bolsa Família position yields PT in 90.9% of coded recommendations.
- A left abortion position is dispersed: PT 45.9%, PSOL 33.9%, and PL 13.8%.
- A right privatization position yields NOVO in 26.3% and PL in 62.6%.
- A right security position yields PL in 94.0%.

The topic matters: some questions spread recommendations across several parties, while others concentrate them on one major party.

Module A, all usable systems, among answers where the script identifies a party.

We compare chatbots using the same voter descriptions

1. For each system and prompt cell, identify the most common answer across repeated calls.
2. Match two systems on prompt cells observed for both.
3. Count how often their most common answers agree.

Broad agreement

- same response type
- same ideological side

Exact agreement

- same party
- or same non-party outcome

The reported mean summarizes 45 pairs among systems with at least 354 Module C cells.

The same voter description usually gets different answers across chatbots

37.7%

mean agreement on response side/type

33.9%

mean agreement on exact party/type

- Each pair compares the same assigned voter descriptions.
- The comparison covers 10 systems with at least 354 of 384 Module C cells.
- The most common answer is taken across repeated calls within each cell.

Forty-five system pairs. GPT-5 remains in the all-answer analysis but has only 52 Module C cells.

Grok and Mistral are the closest pair; Gemini Pro is an outlier

System pair	Same side/type	Same exact party/type
Grok vs Mistral	75.8%	59.4%
Gemini Flash vs Mistral	72.4%	62.5%
Gemini Flash vs Grok	63.3%	47.1%
Gemini Pro vs Mistral	4.2%	4.9%
Gemini Pro vs Sabiá 4	4.2%	4.4%
Gemini Pro vs Grok	2.3%	2.3%

Gemini Pro rarely names a party, so its most common answer differs sharply from chatbots that usually recommend one.

A rich profile is a fictional voter described with many details

- Each profile combines occupation, income, religion, geography, political history, values, and several policy positions.
- The nine profiles adapt voter groups described in *Brasil no Espelho*.
- Examples include a left activist, a low-income voter, a social liberal, a solo entrepreneur, a Christian conservative, an agribusiness voter, and a business owner.

The purpose is to test how chatbots respond to realistic combinations of voter characteristics, not isolated one-line signals.

The profiles are fictional. They do not estimate how real members of these groups would be treated.

We compare whole profiles; no single characteristic is isolated

What changes together


- occupation and income
- identity and past vote
- religion and values
- several policy positions

What we cannot separate

- the effect of race alone
- the effect of class alone
- the effect of religion alone
- the effect of occupation alone

A difference between two profiles belongs to the complete descriptions, not to any one characteristic.

The profiles line up with how their source groups voted in 2022



Nine fictional profiles; average position of the named party across chatbots that name a party on the ideology scale.

Profiles from more pro-Lula groups receive more left-wing recommendations

-0.88

rank correlation As Lula's 2022 vote share rises in the source group, the average named party moves left.

- Left activist, low-income, progressive, and social-liberal profiles appear on the left.
- Agribusiness, Christian conservative, business-owner, and far-right profiles appear on the right.
- The ordering is plausible, but it does not show that any recommendation is correct or fair.

Nine complete profiles; only answers naming a party on the ideology scale enter the correlation.

Some profiles are much more likely to receive a party recommendation

Complete profile	Names party	Refusal/disclaimer
Left activist	86.0%	4.5%
Social liberal	76.8%	11.6%
Far right	63.6%	17.0%
Low-income state-dependent	56.0%	25.0%
Christian conservative	49.2%	18.5%

A follow-up experiment should change one profile characteristic at a time to identify which detail drives the answer.

The preta description reduces refusal by 19.6 percentage points

Description, compared with branca	Refusal/disclaimer	Party named
Parda	-12.1 pp	+14.6 pp
Preta	-19.6 pp	+16.7 pp

In this experiment, both racial descriptions reduce refusal and increase party naming relative to branca. All four confidence intervals exclude zero. Human review of the automated response classification remains important.

All usable Part D answers; comparisons within chatbot, profile step, ballot, and request format.

The same chatbot often changes its answer across identical calls

System	All repeats identical	System	All repeats identical
Claude Opus	14.1%	Grok	48.4%
Claude Sonnet	19.3%	Mistral	60.9%
Gemini Flash	29.9%	Gemini Pro	84.4%
GPT-4o	26.6%	DeepSeek	52.5%
Llama	25.8%	GPT-5	94.2%
Sabiá 4	9.6%		

The outcome combines exact party, refusal, and other response. High agreement can therefore reflect consistently declining or avoiding a party recommendation.

Share of observed Module C cells with identical outcomes across all available repetitions.

Three findings: the chatbot, the voter description, and repeated calls all matter

- **Different chatbots:** two systems give the same most common side or answer type in only 37.7% of matched comparisons.
- **Different wording:** changing identity from left to right makes a party 4.8 percentage points less likely to be named.
- **Same prompt, new call:** many chatbots change the party or answer type across identical requests.

The result is not one shared left-right bias. Each chatbot applies its own, partly inconsistent way of deciding whether and what to recommend.

What the experiment can tell us

- Because identity and policy wording were randomly assigned, average answer differences can be attributed to those wording changes.
- We can compare the named chatbot versions using the same prompts.
- We can describe how complete rich profiles are treated.
- We can measure how often the same chatbot changes its answer across identical calls.
- These conclusions apply to the prompts, systems, and collection period studied here.

What the experiment cannot tell us

- Whether chatbot advice changes how real voters vote
- Whether a recommended party truly matches the stated policy
- Why one vendor or model family behaves differently
- Which single characteristic drives a result for a rich profile
- Whether the same pattern will hold after an interface or model update

Why this matters: voting advice changes across chatbots and across calls

1. Some chatbots routinely name parties; others usually refuse or avoid doing so.
2. Identity and policy information change both whether a party is named and which side is recommended.
3. The same voter description often produces different answers across chatbots.
4. Even the same chatbot may change its answer when the identical request is repeated.

For voters, the platform itself becomes an unobserved part of the advice they receive.

Next: test whether the advice is substantively correct

- Have human coders verify refusals, vague answers, parties, and candidates.
- Link recommended parties and candidates to documented issue positions.
- Compare each recommendation with the policy position stated by the voter.
- Model the two decisions together: whether to name a party and which party to name.
- Repeat the study within a fixed model-version window to measure change over time.

This pilot shows that policy information changes advice. The next question is whether the resulting advice is accurate.

References

-  Bolognesi, Ribeiro, and Codato. 2023. “Classificação ideológica dos partidos políticos brasileiros.” *Dados* 66(2).
-  Miyazaki and Hall. 2026. “Why Do AI Models Tell Left-Wing Voters to Support the Communist Party?” Working paper.
-  Neto. 2024. *Brasil no Espelho*. Quaest, chapter 8.
-  Zucco and Power. 2024. Brazilian Legislative Surveys, waves 1–9. *Latin American Politics and Society* 66(1), appendix C.

All empirical results are authors’ calculations from the July 2026 São Paulo pilot.